

Final Project Data Model(CPE232)

Project : Mercedes Benz Price Prediction

Member Group

- Ms. Sarinyapak Namwong 64070503403
- Mr. Jumponpatha Chaimongkonrojna 64070503408
- Mr. Phanasorn Srisayam 64070503436
- Mr. Warut Wannasert 64070503448
- Mr. Veerapat Toonhirunkorn 64070503474

SUBMITTED TO Prof. SANTITHAM PROM-ON

Import library & Dataset

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

from google.colab import drive
drive.mount('/content/drive')
mercedes = pd.read_csv('/content/drive/MyDrive/Final Project _ Data Model/mercedes-price.csv')

Mounted at /content/drive
```

Preview Data

mercedes

	model	year	price	transmission	mileage	fuelType	tax	mpg	engineSize	
0	SLK	2005	5200	Automatic	63000	Petrol	325	32.1	1.8	
1	S Class	2017	34948	Automatic	27000	Hybrid	20	61.4	2.1	
2	SL CLASS	2016	49948	Automatic	6200	Petrol	555	28.0	5.5	
3	G Class	2016	61948	Automatic	16000	Petrol	325	30.4	4.0	
4	G Class	2016	73948	Automatic	4000	Petrol	325	30.1	4.0	
...	
13114	C Class	2020	35999	Automatic	500	Diesel	145	55.4	2.0	
13115	B Class	2020	24699	Automatic	2500	Diesel	145	55.4	2.0	
13116	GLC Class	2019	30999	Automatic	11612	Diesel	145	41.5	2.1	
13117	CLS Class	2019	37990	Automatic	2426	Diesel	145	45.6	2.0	
13118	S Class	2019	54999	Automatic	2075	Diesel	145	52.3	2.9	

13119 rows x 9 columns

Cleaning Data

Check Data type info

```
mercedes.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13119 entries, 0 to 13118
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   model                13119 non-null object
1   year                 13119 non-null int64
2   price                13119 non-null int64
3   transmission         13119 non-null object
4   mileage              13119 non-null int64
5   fuelType             13119 non-null object
6   tax                  13119 non-null int64
7   mpg                  13119 non-null float64
8   engineSize           13119 non-null float64
dtypes: float64(2), int64(4), object(3)
memory usage: 922.6+ KB
```

Check Null value

```
mercedes.isnull().any()

model          False
year            False
price           False
transmission    False
mileage         False
fuelType        False
tax             False
mpg             False
engineSize      False
dtype: bool
```

Shape function

```
mercedes.shape

(13119, 9)
```

Preprocess data

mercedes.head()

	model	year	price	transmission	mileage	fuelType	tax	mpg	engineSize
0	SLK	2005	5200	Automatic	63000	Petrol	325	32.1	1.8
1	S Class	2017	34948	Automatic	27000	Hybrid	20	61.4	2.1
2	SL CLASS	2016	49948	Automatic	6200	Petrol	555	28.0	5.5
3	G Class	2016	61948	Automatic	16000	Petrol	325	30.4	4.0
4	G Class	2016	73948	Automatic	4000	Petrol	325	30.1	4.0

Checking frequency of 'transmission'

```
pd.crosstab(mercedes['transmission'], columns='Count')
```

col_0	Count
transmission	
Automatic	4825
Manual	1444
Other	2
Semi-Auto	6848

- Note that, from column of transmission which contain Automatic, Manual, Semi-Auto and Other would see a 'Other' column were have only 2 data which could cause outlier.

Cleansing data by deleting transmission type 'Other'

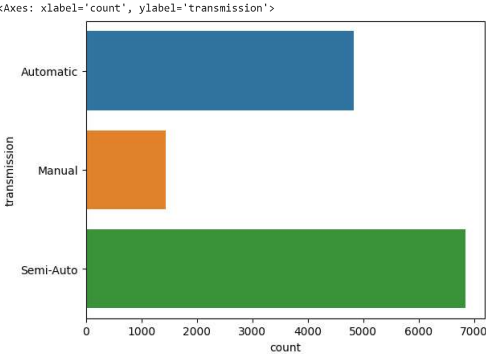
```
cleansing1 = mercedes[mercedes['transmission'] != 'Other'].dropna()
cleansing1
```

	model	year	price	transmission	mileage	fuelType	tax	mpg	engineSize
0	SLK	2005	5200	Automatic	63000	Petrol	325	32.1	1.8
1	S Class	2017	34948	Automatic	27000	Hybrid	20	61.4	2.1
2	SL CLASS	2016	49948	Automatic	6200	Petrol	555	28.0	5.5
3	G Class	2016	61948	Automatic	16000	Petrol	325	30.4	4.0
4	G Class	2016	73948	Automatic	4000	Petrol	325	30.1	4.0
...
13114	C Class	2020	35999	Automatic	500	Diesel	145	55.4	2.0
13115	B Class	2020	24699	Automatic	2500	Diesel	145	55.4	2.0
13116	GLC Class	2019	30999	Automatic	11612	Diesel	145	41.5	2.1
13117	CLS Class	2019	37990	Automatic	2426	Diesel	145	45.5	2.0
13118	S Class	2019	54999	Automatic	2075	Diesel	145	52.3	2.9

13117 rows x 9 columns

Showing 'transmission' data which had cleansed

```
sns.countplot(data=cleansing1, y='transmission')
```



Checking frequency of 'fuelType'

```
pd.crosstab(mercedes['fuelType'], columns='Count')
```

col_0	Count
fuelType	
Diesel	9187
Hybrid	173
Other	7
Petrol	3752

- Note that, from column of fuelType which contain Automatic, Manual, Semi-Auto and Other would see a 'Other' column were have only 2 data which could cause outlier.

Cleansing data by deleting fuelType type 'Other'

```
cleansing2 = cleansing1[cleansing1['fuelType'] != 'Other'].dropna()
cleansing2
```

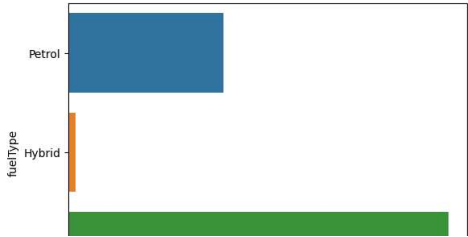
	model	year	price	transmission	mileage	fuelType	tax	mpg	engineSize
0	SLK	2005	5200	Automatic	63000	Petrol	325	32.1	1.8
1	S Class	2017	34948	Automatic	27000	Hybrid	20	61.4	2.1
2	SL CLASS	2016	49948	Automatic	6200	Petrol	555	28.0	5.5
3	G Class	2016	61948	Automatic	16000	Petrol	325	30.4	4.0
4	G Class	2016	73948	Automatic	4000	Petrol	325	30.1	4.0
...
13114	C Class	2020	35999	Automatic	500	Diesel	145	55.4	2.0
13115	B Class	2020	24699	Automatic	2500	Diesel	145	55.4	2.0
13116	GLC Class	2019	30999	Automatic	11612	Diesel	145	41.5	2.1
13117	CLS Class	2019	37990	Automatic	2426	Diesel	145	45.5	2.0
13118	S Class	2019	54999	Automatic	2075	Diesel	145	52.3	2.9

13111 rows x 9 columns

Showing 'transmission' data which had cleansed

```
sns.countplot(data=cleansing2, y='fuelType')
```

<Axes: xlabel='count', ylabel='fuelType'>



Describing Data

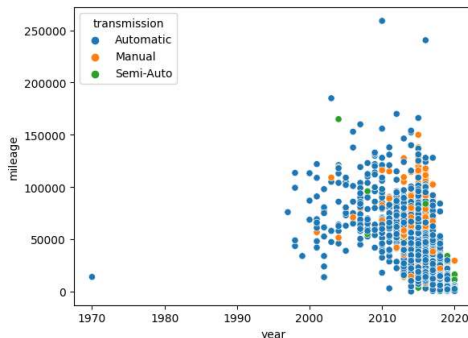
```
mercedes2 = cleansing2.loc[:, 'price': 'engineSize']
mercedes2.describe()
```

	price	mileage	tax	mpg	engineSize
count	13111.000000	13111.000000	13111.000000	13111.000000	13111.000000
mean	24700.935016	21935.920067	130.020212	55.117443	2.071642
std	11843.607522	21169.443845	65.230548	15.094034	0.572256
min	650.000000	1.000000	0.000000	1.100000	0.000000
25%	17450.000000	6090.000000	125.000000	45.600000	1.800000
50%	22480.000000	15181.000000	145.000000	56.500000	2.000000
75%	28980.000000	31733.000000	145.000000	64.200000	2.100000
max	159999.000000	259000.000000	580.000000	217.300000	6.200000

Showing 'fuelType' data which had cleansed

```
sns.scatterplot(data=cleansing2, x="year", y = "mileage", hue="transmission")
```

<Axes: xlabel='year', ylabel='mileage'>



- Find Outlier, so we should cleansing again by delete 1970 data.

```
mercedes2 = cleansing2[cleansing2['year'] != 1970].dropna()
mercedes2
```

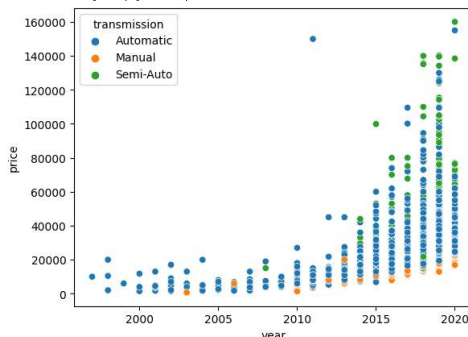
	model	year	price	transmission	mileage	fuelType	tax	mpg	engineSize
0	SLK	2005	5200	Automatic	63000	Petrol	325	32.1	1,8
1	S Class	2017	34948	Automatic	27000	Hybrid	20	61.4	2,1
2	SL CLASS	2016	49948	Automatic	6200	Petrol	555	28,0	5,5
3	G Class	2016	61948	Automatic	16000	Petrol	325	30.4	4,0
4	G Class	2016	73948	Automatic	4000	Petrol	325	30.1	4,0
...
13114	C Class	2020	35999	Automatic	500	Diesel	145	55.4	2,0
13115	B Class	2020	24699	Automatic	2500	Diesel	145	55.4	2,0
13116	GLC Class	2019	30999	Automatic	11612	Diesel	145	41.5	2,1
13117	CLS Class	2019	37990	Automatic	2426	Diesel	145	45,8	2,0
13118	S Class	2019	54999	Automatic	2075	Diesel	145	52,3	2,9

13110 rows × 9 columns

Number of sold car model each year

```
sns.scatterplot(data=mercedes2, x="year", y = "price", hue="transmission")
```

<Axes: xlabel='year', ylabel='price'>



Describing data

```
mercedes2.describe()
```

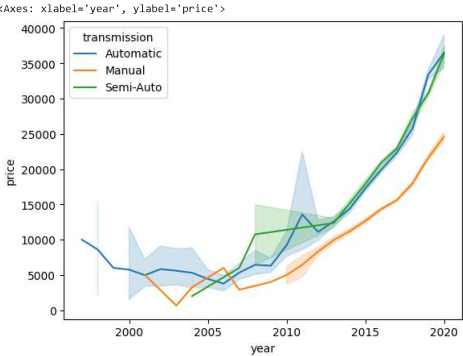
https://colab.research.google.com/drive/1V0OIV_mXx_yd5_3o-x6kMHdWMJsuziO#scrollTo=YMIKpHYAjTb7&printMode=true

3/7

	year	price	mileage	tax	mpg	engineSize
count	13110.000000	13110.000000	13110.000000	13110.000000	13110.000000	13110.000000
mean	2017.300381	24700.912281	21936.525400	130.006865	55.118658	2.071800
std	2.186392	11844.058963	21170.137793	65.215130	15.093970	0.571991
min	1997.000000	650.000000	1.000000	0.000000	1.100000	0.000000
25%	2016.000000	17450.000000	6087.000000	125.000000	45.600000	1.800000
50%	2018.000000	22480.000000	15181.500000	145.000000	56.500000	2.000000
...


Show a price each transmission and each year

```
sns.lineplot(data=mercedes2, x="year", y="price", hue="transmission")
```



Boxplot of price


```
mercedes2.boxplot(column='price', color='orange')
```



Histogram of mercedes model

```
ax = sns.histplot(data=mercedes2, x="model", kde=True, color='orange')
ax.set_xticklabels(ax.get_xticklabels(), rotation=90)
plt.show()

<ipython-input-22-c7d06b28821e:2: UserWarning: FixedFormatter should only be used together with FixedLocator
ax.set_xticklabels(ax.get_xticklabels(), rotation=90)>
```



Histogram of price

```
sns.histplot(data=mercedes2, x="price", kde=True, color='orange')
```

Relation graph between model and price.


```
model_mileage = mercedes2.groupby('model')['price'].mean().sort_values()
plt.figure(figsize=(15,5))
pal = sns.color_palette("Oranges", len(model_mileage))
sns.lineplot(x=model_mileage.index, y=model_mileage, palette=pal)

plt.xticks(rotation=70)

plt.title('Model vs Price graph')

plt.savefig("model_mileage.png")

<ipython-input-24-e940c74c0a7b>:4: UserWarning: Ignoring 'palette' because no 'hue' variable has been assigned.
sns.lineplot(x=model_mileage.index, y=model_mileage, palette=pal)
```



Correlation

```
mercedes2.corr()

<ipython-input-28-07e1643968ce>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only
mercedes2.corr()
```

	year	price	mileage	tax	mpg	engineSize
year	1.000000	0.529856	-0.751821	0.016777	-0.099067	-0.150917
price	0.529856	1.000000	-0.537120	0.268688	-0.442304	0.516591
mileage	-0.751821	-0.537120	1.000000	-0.159173	0.202571	0.063691
tax	0.016777	0.268688	-0.159173	1.000000	-0.515186	0.339955
mpg	-0.099067	-0.442304	0.202571	-0.515186	1.000000	-0.343336
engineSize	-0.150917	0.516591	0.063691	0.339955	-0.343336	1.000000

Correlation Heatmap

```
#sns.set(font_scale=0.8)
sns.heatmap(mercedes2.corr(), annot=True)
plt.show()

<ipython-input-30-01cefc933d82>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only
sns.heatmap(mercedes2.corr(), annot=True)
```



Modeling

```
from statsmodels.formula.api import ols

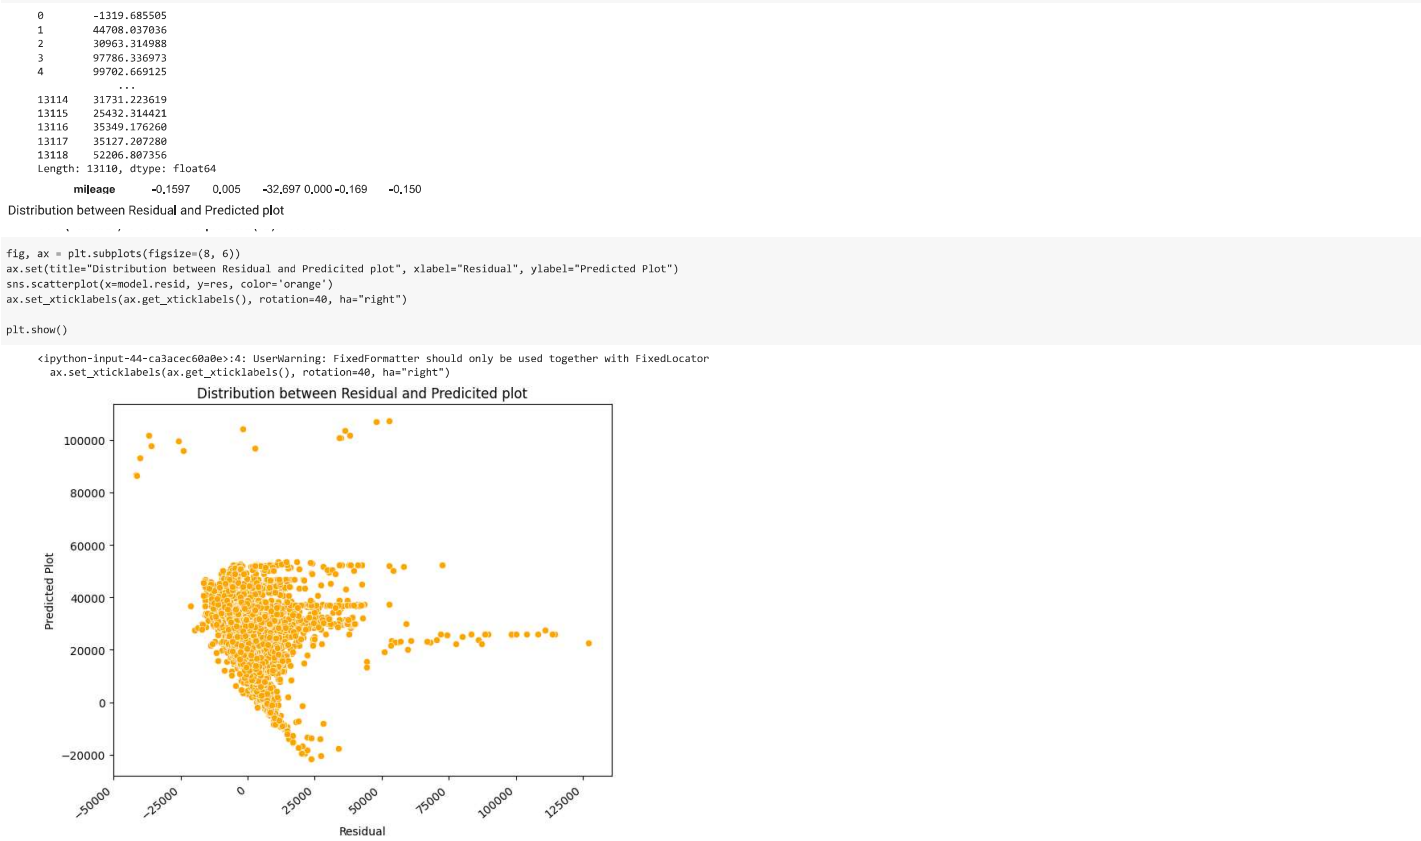
model = ols('price ~ year + model + mileage',mercedes2, missing='drop').fit()
model.summary()
```

```
OLS Regression Results
Dep. Variable: price R-squared: 0.584
Model: OLS Adj. R-squared: 0.584
Method: Least Squares F-statistic: 657.1
Date: Thu, 18 May 2023 Prob (F-statistic): 0.00
Time: 05:06:05 Log-Likelihood: -1.3581e+05
No. Observations: 13110 AIC: 2.717e+05
Df Residuals: 13081 BIC: 2.719e+05
Df Model: 28
Covariance Type: nonrobust

coef std err t P>|t| [0.025 0.975]
Intercept -3.526e+06 9.99e+04 -35.280 0.000 -3.72e+06 -3.33e+06
model[T. B Class] -1934.0683 349.020 -5.541 0.000 -2618.199 -1249.938
model[T. C Class] 4045.4522 196.084 20.631 0.000 3661.099 4429.805
model[T. CL Class] 2958.9588 370.586 7.985 0.000 2232.555 3685.362
model[T. CLA Class] 3071.7647 838.476 3.664 0.000 1428.231 4715.299
model[T. CLC Class] 4374.3458 4423.257 0.989 0.323 -4295.880 1.3e+04
model[T. CLK] 2.173e+04 2929.760 7.418 0.000 1.6e+04 2.75e+04
model[T. CLS Class] 9508.2015 519.746 18.294 0.000 8489.424 1.05e+04
model[T. E Class] 6387.2715 229.843 27.790 0.000 5936.746 6837.797
model[T. G Class] 7.861e+04 1979.464 40.219 0.000 7.57e+04 8.35e+04
model[T. GL Class] 6627.7082 713.375 9.291 0.000 5229.390 8026.026
model[T. GLA Class] 1407.0666 303.226 4.640 0.000 812.699 2001.434
model[T. GLB Class] 1.013e+04 1762.041 5.748 0.000 6673.581 1.36e+04
model[T. GLC Class] 1.12e+04 290.095 38.598 0.000 1.06e+04 1.18e+04
model[T. GLE Class] 1.922e+04 387.228 49.642 0.000 1.85e+04 2e+04

res = model.predict(mercedes2)
res

0 -1319.685505
1 44708.037036
2 39963.314988
3 97786.336973
4 99702.669125
...
13114 31731.223619
13115 25432.314421
13116 35349.176260
13117 35127.207280
13118 52206.807356
Length: 13110, dtype: float64
mileage -0.1597 0.005 -32.697 0.000 -0.169 -0.150
```



```
Accuracy

mercedes_num = mercedes2.select_dtypes('number')
mercedes_object = mercedes2.select_dtypes('object')

mercedes_object1 = pd.get_dummies(mercedes_object, drop_first=True, columns=(['transmission','fuelType','model']))
#mercedes_object1 = pd.get_dummies(mercedes_object, drop_first=True)

mercedes_prep = pd.concat([mercedes_num,mercedes_object1], axis=1)
mercedes_prep
```

	year	price	mileage	tax	mpg	engineSize	transmission_Manual	transmission_Semi-Auto	fuelType_Hybrid	fuelType_Petrol	...	model_R Class	model_S Class	model_SL CLASS	model_SLK	model_V Class	model_X-CLASS	model_180	model_200	model_220	model_230
0	2005	5200	63000	325	32.1	1.8	0	0	0	1	...	0	0	0	1	0	0	0	0	0	0
1	2017	34948	27000	20	61.4	2.1	0	0	1	0	...	0	1	0	0	0	0	0	0	0	0
2	2016	49948	6200	555	28.0	5.5	0	0	0	1	...	0	0	1	0	0	0	0	0	0	0
3	2016	61948	16000	325	30.4	4.0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	0
4	2016	73948	4000	325	30.1	4.0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	0
...
13114	2020	35999	500	145	55.4	2.0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
13115	2020	24699	2500	145	55.4	2.0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
13116	2019	30999	11612	145	41.5	2.1	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
13117	2019	37990	2426	145	45.6	2.0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
13118	2019	54999	2075	145	52.3	2.9	0	0	0	0	...	0	1	0	0	0	0	0	0	0	0

13110 rows x 36 columns

```
from sklearn.model_selection import train_test_split

mercedes_train, mercedes_test = train_test_split(mercedes_prep, train_size = 0.7)

mercedes_train.shape
```

```
(9177, 36)

mercedes_test.shape

(3933, 36)

from sklearn.linear_model import LinearRegression

tree = LinearRegression()
tree.fit(mercedes_train.drop(columns='price'),mercedes_train['price'])
accuracy = tree.score(mercedes_test.drop(columns='price'),mercedes_test['price'])
print('Accuracy:',accuracy)

Accuracy: 0.7636114295699339
```